

Fine-Tuning Florence-2 on a License Plate Dataset with EasyOCR

Project Overview

This project fine-tunes Microsoft's **Florence-2** model on a **license plate dataset** using **EasyOCR**. The goal is to improve the model's accuracy in detecting and recognizing license plates from images.

Dependencies Installation

The notebook begins by installing the required Python libraries:

```
!pip install einops sklearn python-Levenshtein datasets timm jiwer
```

These libraries are essential for:

- **Image processing** ([einops](#), [timm](#))
- **Levenshtein distance for OCR accuracy** ([python-Levenshtein](#))
- **Dataset handling** ([datasets](#))
- **Evaluation metrics** ([jiwer](#))

Loading Florence-2 Model & Processor

```
from transformers import AutoModelForCausalLM, AutoProcessor
```

```
import torch
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
print(device)
```

```
model = AutoModelForCausalLM.from_pretrained(
```

```
    "microsoft/Florence-2-base-ft", trust_remote_code=True, revision='refs/pr/6'
```

```
).to(device)
```

```
processor = AutoProcessor.from_pretrained(
    "microsoft/Florence-2-base-ft", trust_remote_code=True, revision='refs/pr/6'
)

torch.cuda.empty_cache()
```

- **Loads the Florence-2 model**, optimized for image-to-text tasks.
- **Detects GPU availability** and loads the model onto it for faster processing.
- **Clears CUDA cache** to free up memory.



Loading the License Plate Dataset

```
import datasets
```

```
dataset = datasets.load_dataset("keremberke/license-plate-object-detection", 'mini')
```

```
dataset
```

This dataset contains:

- Images of license plates 📷
- Corresponding text labels of detected plates abc
- Bounding box information 📏



Inspecting Dataset Samples

```
for idx in range(len(dataset["train"])):
```

```
    sample = dataset["train"][idx]
```

```
    print(f"🔍 Image ID: {sample['image_id']}")
```

```
    print(f"📷 Image: {sample['image']}")
```

```
    print(f"✅ Ground Truth Label: {sample['objects']}")
```

- **Loops through training samples** and prints:
 - Image ID

- Image data
- Ground truth label (actual text on license plate)

Splitting Dataset into Training, Testing & Validation Sets

```
ds_train = dataset["train"]  
  
ds_test = dataset["test"]  
  
ds_val = dataset["validation"]
```

- **Divides data** into training, validation, and test sets.

Running Inference on the Model

Function to run the model on an example

```
def run_example(task_prompt, text_input, image):
```

```
    prompt = task_prompt + text_input
```

```
    # Ensure the image is in RGB mode
```

```
    if image.mode != "RGB":
```

```
        image = image.convert("RGB")
```

```
    inputs = processor(text=prompt, images=image, return_tensors="pt").to(device)
```

```
    generated_ids = model.generate(  
        input_ids=inputs["input_ids"],
```

```
        pixel_values=inputs["pixel_values"],  
  
        max_new_tokens=1024,
```

```
        num_beams=3
```

```
    )
```

```
    generated_text = processor.batch_decode(generated_ids, skip_special_tokens=False)[0]
```

```
parsed_answer = processor.post_process_generation(generated_text, task=task_prompt,
image_size=(image.width, image.height))
```

```
return parsed_answer
```

- **Takes an image and runs the model inference.**
- **Extracts the license plate text** based on a given question.

Testing the Model on Sample Images

```
# Test on the first 3 samples from the train dataset
```

```
for idx in range(3):
```

```
    sample = dataset["train"][idx]
```

```
    image = sample["image"]
```

```
    print(run_example("DocVQA", "What is the license plate number?", image))
```

```
    display(image)
```

- **Runs inference on three sample images.**
- **Displays images and predicted license plate numbers.**

Evaluating OCR Accuracy

```
from jiwer import wer
```

```
def calculate_ocr_accuracy(predictions, ground_truths):
```

```
    errors = [wer(gt, pred) for gt, pred in zip(ground_truths, predictions)]
```

```
    return 1 - sum(errors) / len(errors)
```

- Uses **Word Error Rate (WER)** to measure OCR accuracy.
- Compares predicted license plate text with ground truth labels.

Frequently Asked Questions (FAQs)

❶ Why is the OCR result sometimes incorrect?

OCR errors can occur due to:

- Blurry or low-resolution images 📷
- Poor lighting conditions 🌑
- Unusual fonts or characters 📄 abc

❷ How can OCR accuracy be improved?

- Increase dataset size 📊
- Use higher-quality images 📷
- Fine-tune the model further with more epochs 🏃

❸ Why does the model fail on some images but not others?

Different conditions such as angle, lighting, and occlusions affect the recognition performance. Proper dataset augmentation can help!

❹ How is OCR accuracy measured in this project?

We use **Word Error Rate (WER)** and **Levenshtein distance**, which compare the predicted text with the ground truth label.

❺ Can this model be used in real-time applications?

Yes! With GPU acceleration, inference time is reduced, making it suitable for **real-time license plate recognition** in surveillance and traffic systems 🚦.

🔑 **Client Takeaway:** This documentation provides a step-by-step breakdown of the model training, dataset processing, and OCR evaluation for license plate recognition using **Microsoft Florence-2 and EasyOCR**.